

## › CNN vs MLP on MNIST

### IFT6135-H2019 Assignment 1

• Akila Daniel Jeeson • Myriam Laiymani • Marc-André Piché • Khalil Slimi

#### › Initialising Colab drive

↳ 3 cellules masquées

#### › Librairies

↳ 1 cellule masquée

#### › Dataset

```
[ ] def showRandSample(dataset, n_row=5, n_col=5):
    nb_img = n_row * n_col

    sample_loader = DataLoader(
        dataset,
        batch_size=nb_img,
        shuffle=True,
        num_workers=1,
        pin_memory=cuda)

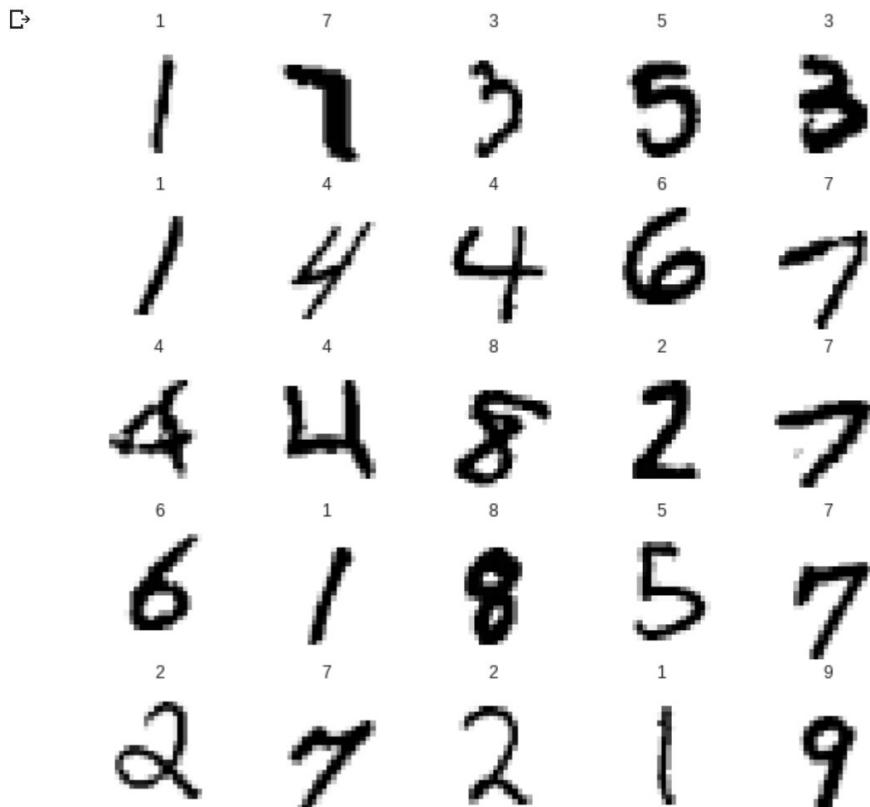
    images, labels = next(iter(sample_loader))
    images = images.squeeze()

    fig = plt.figure(figsize=(n_row*2, n_col*2))

    for i in range(1, nb_img +1):
        img = images[i-1]
        ax = fig.add_subplot(n_row, n_col, i)
        ax.set_title(str(labels[i-1].item()))
        plt.axis('off')
        plt.imshow(img,cmap='gray_r')

    plt.show()
```

```
[ ] showRandSample(train_data)
```



\*The best way to view this document is to simply close it and open the source 6135-TP1problem\_2.ipynb in Colab



```

# Fully connected layer
self.fc1 = nn.Linear(32*4*4, 768)
self.fc2 = nn.Linear(768, 10)

def forward(self, x):
    x = F.relu(self.conv1(x))
    x = F.max_pool2d(x, 7)

    x = x.view(x.size(0), -1) # flatten

    x = F.relu(self.fc1(x))
    x = self.fc2(x)

    return x

class CNN_2(nn.Module):

    def __init__(self):
        super().__init__()
        # Conv block
        self.conv1 = nn.Conv2d(1, 87, 3, padding=2)
        self.conv2 = nn.Conv2d(87, 512, 3, padding=1)

        # Fully connected layer
        self.fc1 = nn.Linear(512*1*1, 768)
        self.fc2 = nn.Linear(768, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.max_pool2d(x, 3)
        x = F.relu(self.conv2(x))
        x = F.max_pool2d(x, 10)

        x = x.view(x.size(0), -1) # flatten

        x = F.relu(self.fc1(x))
        x = self.fc2(x)

        return x

class CNN_3(nn.Module):

    def __init__(self):
        super().__init__()
        # Conv block
        self.conv1 = nn.Conv2d(1, 64, 3, padding=3)
        self.conv2 = nn.Conv2d(64, 77, 3, padding=1)
        self.conv3 = nn.Conv2d(77, 512, 3, padding=1)

        # Fully connected layer
        self.fc1 = nn.Linear(512*1*1, 768)
        self.fc2 = nn.Linear(768, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.max_pool2d(x, 2)
        x = F.relu(self.conv2(x))
        x = F.max_pool2d(x, 4)
        x = F.relu(self.conv3(x))
        x = F.max_pool2d(x, 4)

        x = x.view(x.size(0), -1) # flatten

        x = F.relu(self.fc1(x))
        x = self.fc2(x)

        return x

```

#### ▼ Model verification and number of parameters :

```

#Number of parameters
from torchsummary import summary

model = CNN_1().to(device)

print("CNN_1 : ",sum(p.numel() for p in model.parameters() if p.requires_grad))
summary(model, (1, 28, 28))

model = CNN_2().to(device)

print("\nCNN_2 : ",sum(p.numel() for p in model.parameters() if p.requires_grad))
summary(model, (1, 28, 28))

model = CNN_3().to(device)

print("\nCNN_3 : ",sum(p.numel() for p in model.parameters() if p.requires_grad))
summary(model, (1, 28, 28))

model = MLP_2L().to(device)

print("\nMLP_2L : ",sum(p.numel() for p in model.parameters() if p.requires_grad))
summary(model, (1, 28, 28))

```

```
CNN_1 : 401994
-----
          Layer (type)           Output Shape        Param #
=====
            Conv2d-1           [-1, 32, 28, 28]       320
            Linear-2            [-1, 768]         393,984
            Linear-3            [-1, 10]          7,690
=====
Total params: 401,994
Trainable params: 401,994
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.20
Params size (MB): 1.53
Estimated Total Size (MB): 1.73
-----

CNN_2 : 803952
-----
          Layer (type)           Output Shape        Param #
=====
            Conv2d-1           [-1, 87, 30, 30]      870
            Conv2d-2            [-1, 512, 10, 10]    401,408
            Linear-3            [-1, 768]         393,984
            Linear-4            [-1, 10]          7,690
=====
Total params: 803,952
Trainable params: 803,952
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.99
Params size (MB): 3.07
Estimated Total Size (MB): 4.06
-----

CNN_3 : 802071
-----
          Layer (type)           Output Shape        Param #
=====
            Conv2d-1           [-1, 64, 32, 32]      640
            Conv2d-2            [-1, 77, 16, 16]    44,429
            Conv2d-3            [-1, 512, 4, 4]     355,328
            Linear-4            [-1, 768]         393,984
            Linear-5            [-1, 10]          7,690
=====
Total params: 802,071
Trainable params: 802,071
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.72
Params size (MB): 3.06
Estimated Total Size (MB): 3.78
-----

MLP_2L : 803594
-----
          Layer (type)           Output Shape        Param #
=====
            Linear-1           [-1, 512]        401,920
            Linear-2            [-1, 768]         393,984
            Linear-3            [-1, 10]          7,690
=====
Total params: 803,594
Trainable params: 803,594
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.01
Params size (MB): 3.07
Estimated Total Size (MB): 3.08
-----
```

```
[ ] loss_fn = nn.CrossEntropyLoss()
test_loss_fn = nn.CrossEntropyLoss(reduction='sum')

#Results directory
savedir = 'results'
if not os.path.exists(savedir):
    os.makedirs(savedir)

## Saved states are ignored for this expirement

#checkpoint = torch.load('/content/drive/My Drive/Colab Notebooks//6135/Assignment_1.2')
#model.load_state_dict(checkpoint['state_dict'])

[ ] def test(model, loader, dset = 'validation'):
    model.eval()

    test_loss = 0
    correct = 0
    test_size = 0
    with torch.no_grad():
        for inputs, target in loader:
            inputs = inputs.to(device,dtype=torch.float)
            target = target.to(device, dtype=torch.long)

            output = model(inputs)
            test_size += len(inputs)
            test_loss += test_loss_fn(output, target).item()
            pred = output.max(1, keepdim=True)[1]
            correct += pred.eq(target.view_as(pred)).sum().item()

    test_loss /= test_size
    accuracy = correct / test_size
    if dset == "validation":
        print('Validation error: {} / {} ({:.2f}%) \n'.format(
            (test_size-correct), test_size,
            100. * (1-accuracy)))
        print('Validation set: Average loss: {:.4f}, Accuracy: {} / {} ({:.2f}%) \n'.format(
            test_loss, correct, test_size,
            100. * accuracy))
    if dset == "train":
        print('\nTraining error: {} / {} ({:.2f}%) \n'.format(
            (test_size-correct), test_size,
            100. * (1-accuracy)))
    if dset == "test":
        print('\nTest set Accuracy: {} / {} ({:.2f}%) \n'.format(
            correct, test_size,
            100. * accuracy))

    return test_loss, accuracy
```

```
[ ] def train(model,loader, optimizer, epoch, scheduler):
    model.train()

    for batch_idx, (inputs, target) in enumerate(loader):
        inputs = inputs.to(device, dtype=torch.float)
        target = target.to(device, dtype=torch.long)

        optimizer.zero_grad()
        output = model(inputs)
        loss = loss_fn(output, target)
        loss.backward()
        optimizer.step()

        if batch_idx % 100 == 0:
            print(model.__class__.__name__, ' Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch +1, batch_idx* len(inputs), len(loader) *len(inputs) ,
                100. * batch_idx / len(loader), loss.item()))

    scheduler.step()
    return test(model, loader, dset = 'train')
```

## Running the experiment

The experiment being non-determinist, depending on weight initialisation and the draw of the training set, we run the experiment multiple times and take the average results for each corresponding epochs, resetting all weights and sets for each run.

Here : 10 runs of 10 epochs.

```
[ ] runs = 10
epochs = 10
```

```
[ ] import pickle

def weight_reset(m):
    if isinstance(m, nn.Conv2d) or isinstance(m, nn.Linear):
        m.reset_parameters()

def MultipleRun(runs, epochs, model):
    liveLoss = PlotLosses()
    model = model.to(device)
    name = model.__class__.__name__

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.1)
    exp_lr_scheduler = lr_scheduler.StepLR(optimizer, step_size=10, gamma=1.0)

    results = {'name': name, 'loss': [0]*epochs, 'accuracy': [0]*epochs}
    savefile = os.path.join(savedir, results['name']+'.pkl')

    best_net = '/content/drive/My Drive/Colab Notebooks//6135/Assignment_1.2/best_' + name

    for run in range(runs):
        print("\n", " -- Run : ", run+1)
        since = time.time()
        best_accuracy = 0

        #Resetting all weights for new run
        model.apply(weight_reset)

        log_error = {'train_error': [0]*epochs, 'valid_error': [0]*epochs}

        for epoch in range(epochs):
            train_loss, train_acc = train(model, train_loader, optimizer, epoch, exp_lr_scheduler)
            loss, acc = test(model, valid_loader, dset = 'validation')

            if acc > best_accuracy:
                best_accuracy = acc

            #Saving best model parameters
            torch.save({
                'epoch_based0': epoch,
                'state_dict': model.state_dict(),
                'accuracy': acc,
                'loss' : loss,
                'optimizer' : optimizer.state_dict(),
            }, best_net)

            print('new accuracy parameters saved {}'.format(best_accuracy))

        #Updating average results
        results['loss'][epoch] += loss/runs
        results['accuracy'][epoch] += acc/runs
        log_error['train_error'][epoch] += (1-train_acc)*100
        log_error['valid_error'][epoch] += (1-acc)*100

    _, _ = test(model, test_loader, dset = "test")
    time_elapsed = time.time() - since
    print('Training complete in {:.0f}m {:.0f}s'.format(
        time_elapsed // 60, time_elapsed % 60))

    ##plotting errors
    fig, ax1 = plt.subplots(figsize=(10, 4))
    ax1.plot(log_error['train_error'], label = "train")
    ax1.plot(log_error['valid_error'], label = "validation")
    ax1.set_ylabel('error')
    ax1.set_xlabel('epochs')
    plt.legend()
    plt.show()

    with open(savefile, 'wb') as fout:
        pickle.dump(results, fout)
```

## ▼ Running CNN1

```
[ ] MultipleRun(runs, epochs, CNN_1())

→
-- Run : 1
CNN_1 Epoch: 1 [0/45056 (0%)] Loss: 2.294302
CNN_1 Epoch: 1 [12800/45056 (28%)] Loss: 0.848484
CNN_1 Epoch: 1 [25600/45056 (57%)] Loss: 0.412616
CNN_1 Epoch: 1 [38400/45056 (85%)] Loss: 0.264937

Training error: 4098/45000 (9.11%)
Validation error: 1364/15000 (9.09%)

Validation set: Average loss: 0.3017, Accuracy: 13636/15000 (90.91%)

new accuracy parameters saved 0.9090666666666667
CNN_1 Epoch: 2 [0/45056 (0%)] Loss: 0.218321
CNN_1 Epoch: 2 [12800/45056 (28%)] Loss: 0.237038
CNN_1 Epoch: 2 [25600/45056 (57%)] Loss: 0.346456
CNN_1 Epoch: 2 [38400/45056 (85%)] Loss: 0.156660
```









